

ASYNCHRONOUS EXECUTION

LEARNING OBJECTIVES

- Learn about how commands are enqueued asynchronously
- Learn about the different reasons for synchronization
- Learn about the different ways to perform synchronization

ASYNCHRONOUS EXECUTION

- All command submitted to a `queue` are done so asynchronously.
- The functions return immediately and the command is run in a background thread.
- This includes individual commands like `memcpy` and collections of commands derived from a command group.
- This means you have to synchronize with those commands.

SYNCHRONIZATION

There are a number of reasons why you need to synchronize with commands

- Await completion of a kernel function.
- Await the results of a computation.
- Await error conditions which come from a failure to execute any of the commands.

SYNCHRONIZATION WITH KERNEL FUNCTIONS

There are two ways ways to synchronize with kernel functions.

- Calling `wait` on an `event` object returned from enqueueing a kernel function command, either via a command group or a shortcut function.
- Calling `wait` or `wait_and_throw` on the `queue` itself.

SYNCHRONIZING WITH KERNEL FUNCTIONS (BUFFERS/ACCESSORS)

```
buf = sycl::buffer(data, sycl::range{1024});  
  
gpuQueue.submit([&](sycl::handler &cgh){  
    auto acc = sycl::accessor{buf, cgh};  
  
    cgh.parallel_for<kernel_a>(sycl::range{1024},  
        [=](sycl::id<1> idx){  
            acc[idx] = /* some computation */  
        });  
}).wait();
```

- Calling `wait` on an event object returned from enqueueing a command group will wait for the commands from that command group to complete.
- This is how we have synchronized in our examples so far.
- This effectively creates a blocking operations that will complete in place by immediately synchronizing.

SYNCHRONIZING WITH KERNEL FUNCTIONS (BUFFERS/ACCESSORS)

```
buf = sycl::buffer(data, sycl::range{1024});  
  
gpuQueue.submit([&](sycl::handler &cgh){  
    auto acc = sycl::accessor{buf, cgh};  
  
    cgh.parallel_for<kernel_a>(sycl::range{1024},  
        [=](sycl::id<1> idx){  
            acc[idx] = /* some computation */  
        });  
});  
  
gpuQueue.wait();
```

- Calling `wait` or `wait_and_throw` on a queue will wait for all commands enqueued to it to complete.
- Note that command groups do not create commands to copy data back to the host application.

SYNCHRONIZING WITH KERNEL FUNCTIONS (USM)

```
auto devicePtr = usm_wrapper<int>(
    malloc_device<int>(1024, gpuQueue));

gpuQueue.memcpy(devicePtr, data, sizeof(int)).wait();

gpuQueue.parallel_for<kernel_a>(sycl::range{1024},
    [=](sycl::id<1> idx){
        devicePtr[idx] = /* some computation */
    }).wait();
```

- Calling `wait` on an event object returned from functions such as `memcpy` or the queue shortcuts will wait for that specific command to complete.
- Again this is how we have synchronized in our examples so far.

SYNCHRONIZING WITH KERNEL FUNCTIONS (USM)

```
auto devicePtr = usm_wrapper<int>(
    malloc_device<int>(1024, gpuQueue));

gpuQueue.memcpy(devicePtr, data, sizeof(int));

gpuQueue.wait();

gpuQueue.parallel_for<kernel_a>(sycl::range{1024},
    [=](sycl::id<1> idx){
        devicePtr[idx] = /* some computation */
    });

gpuQueue.wait();
```

- Again calling `wait` or `wait_and_throw` on a queue will wait for all commands enqueued to it to complete.
- Note you generally don't want to call `wait` on the queue after every command, instead you want to create dependencies between commands, which we cover in the next lecture.

SYNCHRONIZING WITH DATA

There are multiple ways ways to synchronize with data, but it differs depending on the data management model you are using.

- When using the USM data management model you can synchronize the same way you would for kernel functions, calling `wait` on an event or the queue.
- When using the buffer/access data management model command groups don't automatically copy data back so there are other ways to synchronize with the data.
 - Creating a `host_accessor`.
 - Destroying the `buffer`.

SYNCHRONIZING WITH DATA (USM)

```
gpuQueue.memcpy(data, devicePtr, sizeof(int)).wait();
```

```
gpuQueue.memcpy(data, devicePtr, sizeof(int));  
gpuQueue.wait();
```

- Simply call `wait` on the event returned from `memcpy`.
- Alternatively call `wait` on the queue.

SYNCHRONIZING WITH DATA (BUFFER/ACCESSOR)

```
buf = sycl::buffer(data, sycl::range{1024});

gpuQueue.submit([&](sycl::handler &cgh){
    auto acc = sycl::accessor{buf, cgh};

    cgh.parallel_for<kernel_a>(sycl::range{1024},
        [=](sycl::id<1> idx){
            acc[idx] = /* some computation */
        });
});

{
    auto hostAcc = buf.get_host_access(sycl::write_only);

    hostAcc[index] = /* some computation */
}
```

- A `host_accessor` gives immediate access to the data managed by a `buffer` in the host application.
- This will wait for any kernel functions accessing the `buffer` to complete and then copying the data back to the host.
- It will also block any other accessor accessing a `buffer` until it is destroyed.
- Note that the data

SYNCHRONIZING WITH DATA (BUFFER/ACCESSOR)

```
{
    buf = sycl::buffer(data, sycl::range{1024});

    gpuQueue.submit([&](sycl::handler &cgh){
        auto acc = sycl::accessor{buf, cgh};

        cgh.parallel_for<kernel_a>(sycl::range{1024},
            [=](sycl::id<1> idx){
                acc[idx] = /* some computation */
            });
    });
}
```

- A buffer will also synchronize the data it manages on destruction.
- It will wait for any kernel functions accessing it to complete and copy the data back to the origin address before completing destruction.

SYNCHRONIZING WITH ERRORS

- Errors are handled by a `queue` and any asynchronous errors can be produced during any of the synchronization methods we've looked at.
- The best way to ensure all errors are caught is to synchronize by calling `wait` or `wait_and_throw` on the `queue`.

QUESTIONS

EXERCISE

`Code_Exercises/Exercise_9_Synchronization/source`

Try out the different methods of synchronizing with a kernel function and the resulting data from the computation.

