

COALESCED GLOBAL MEMORY

LEARNING OBJECTIVES


- Learn about coalesced global memory access
- Learn about the performance impact
- Learn about row-major vs column-major
- Learn about SoA vs AoS

COALESCED GLOBAL MEMORY

- Reading from and writing to global memory is generally very expensive.
- It often involves copying data across an off-chip bus.
 - This means you generally want to avoid unnecessary accesses.
- Memory access operations is done in chunks.
 - This means accessing data this is physically close together in memory is more efficient.

COALESCED GLOBAL MEMORY

```
float data[size];
```




COALESCED GLOBAL MEMORY

```
float data[size];
```

```
...
```

```
f(a[globalId]);
```




COALESCED GLOBAL MEMORY

```
float data[size];
```

```
...
```

```
f(a[globalId]);
```




COALESCED GLOBAL MEMORY

```
float data[size];
```

```
...
```

```
f(a[globalId]);
```

100% global access utilisation




COALESCED GLOBAL MEMORY

```
float data[size];
```

```
...
```

```
f(a[globalId * 2]);
```




COALESCED GLOBAL MEMORY

```
float data[size];
```

```
...
```

```
f(a[globalId * 2]);
```

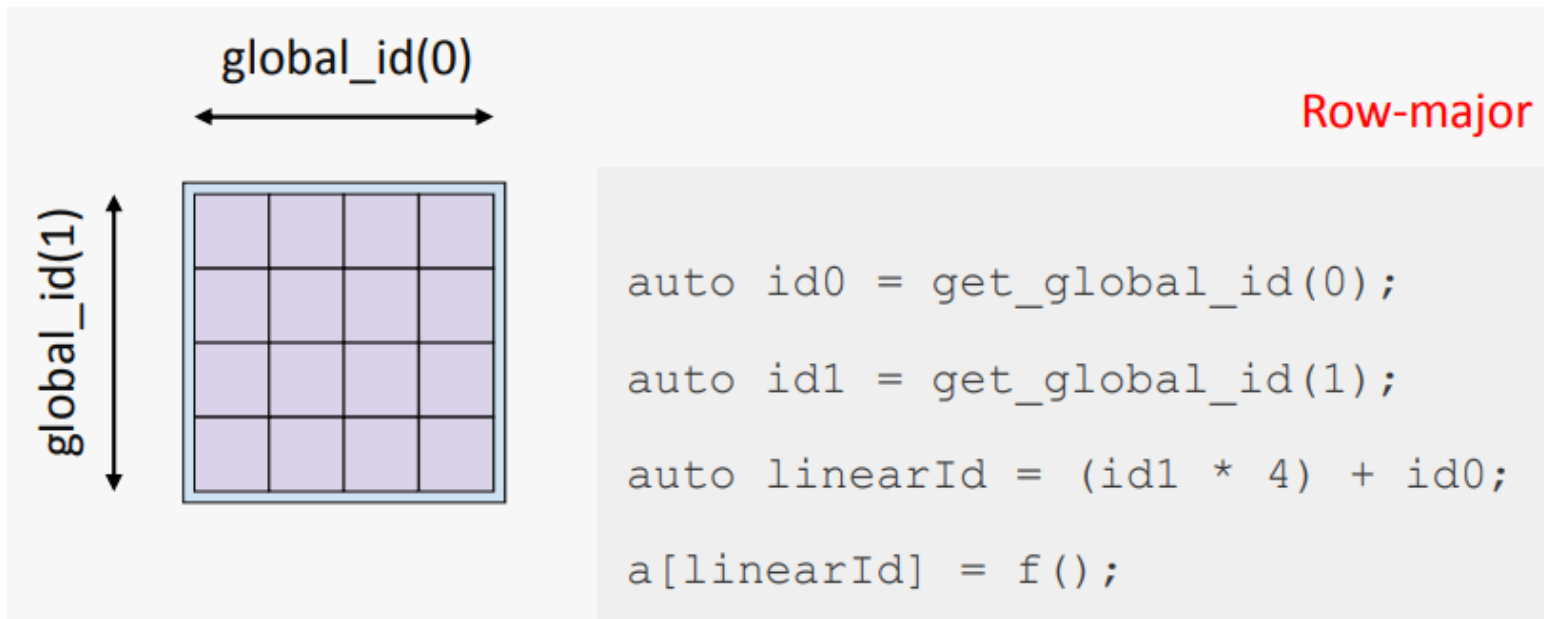
50% global access utilisation




ROW-MAJOR VS COLUMN-MAJOR

- Coalescing global memory access is particularly important when working in multiple dimensions.
- This is because when doing so you have to convert from a position in 2d space to a linear memory space.
- There are two ways to do this; generally referred to as row-major and column-major.

ROW-MAJOR VS COLUMN-MAJOR



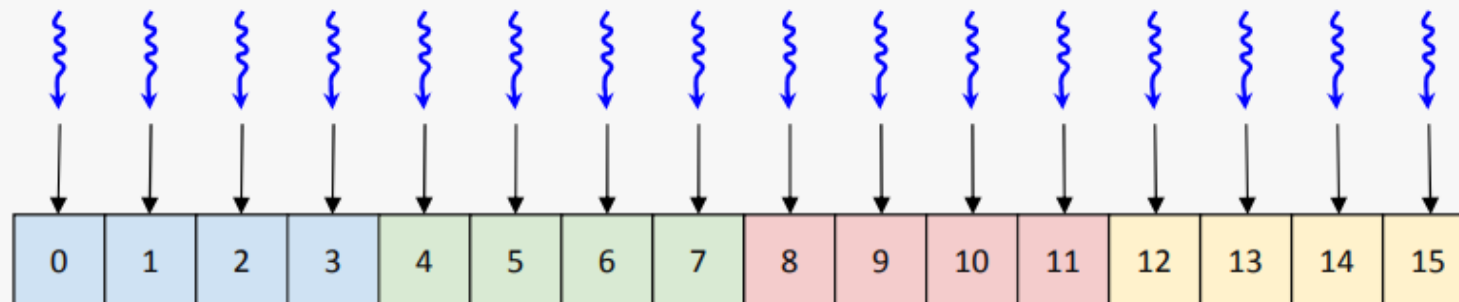
ROW-MAJOR VS COLUMN-MAJOR




Row-major

Row-major

```
auto id0 = get_global_id(0);  
auto id1 = get_global_id(1);  
auto linearId = (id1 * 4) + id0;  
a[linearId] = f();
```




ROW-MAJOR VS COLUMN-MAJOR




Row-major

```
auto id0 = get_global_id(0);
auto id1 = get_global_id(1);
auto linearId = (id1 * 4) + id0;
a[linearId] = f();
```

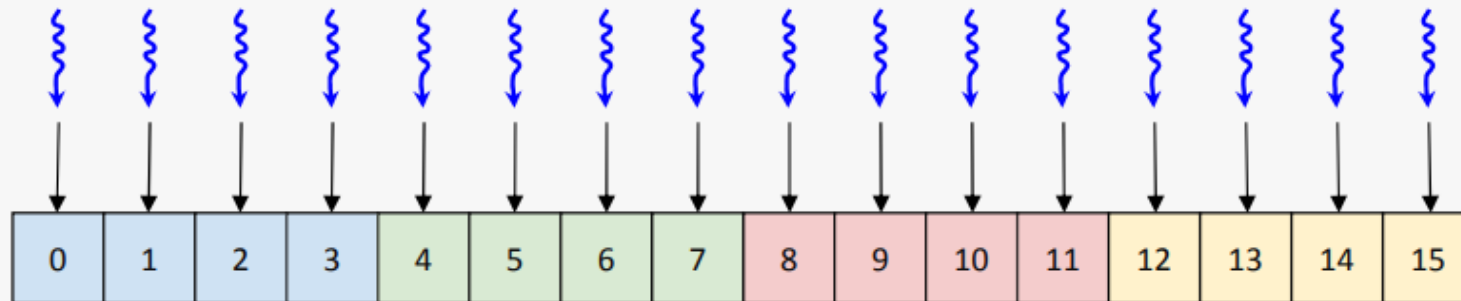


ROW-MAJOR VS COLUMN-MAJOR




Column-major

```
auto id0 = get_global_id(0);  
auto id1 = get_global_id(1);  
auto linearId = (id0 * 4) + id1;  
a[linearId] = f();
```

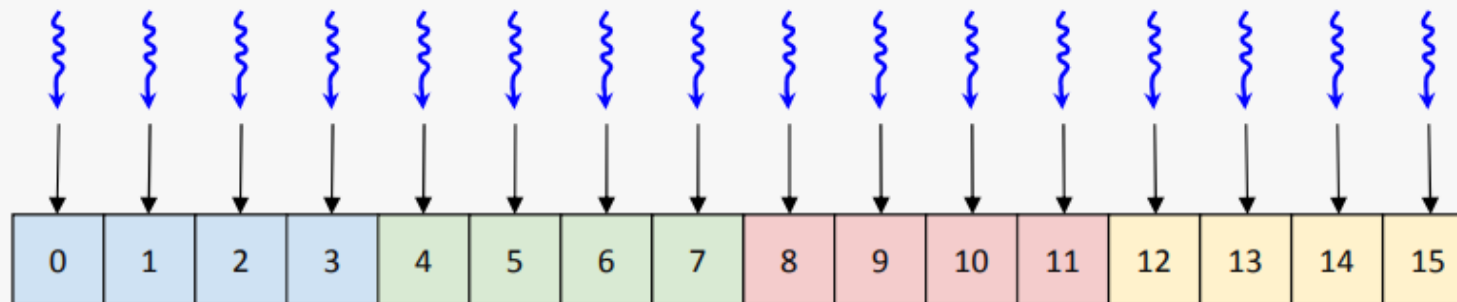


ROW-MAJOR VS COLUMN-MAJOR



Column-major

```
auto id0 = get_global_id(0);  
auto id1 = get_global_id(1);  
auto linearId = (id0 * 4) + id1;  
a[linearId] = f();
```



AOS VS SOA

- Another area this is a factor is when composing data structures.
- It's often instinctive to have struct representing a collection of data and then have an array of this - often referred to as Array of Structs (AoS).
- But for data parallel architectures such as a GPU it's more efficient to have sequential elements of the same type stored contiguously in memory - often referred to as Struct of Arrays (SoA).

AOS VS SOA


```
struct str {
    float f;
    int i;
};
```

```
str data[size];
```

[illegible]


AOS VS SOA

```
struct str {  
    float f;  
    int i;  
};  
  
str data[size];  
  
...  
  
f(a[globalId].f);
```




AOS VS SOA

```
struct str {  
    float f;  
    int i;  
};  
  
str data[size];  
  
...  
  
f(a[globalId].f);
```



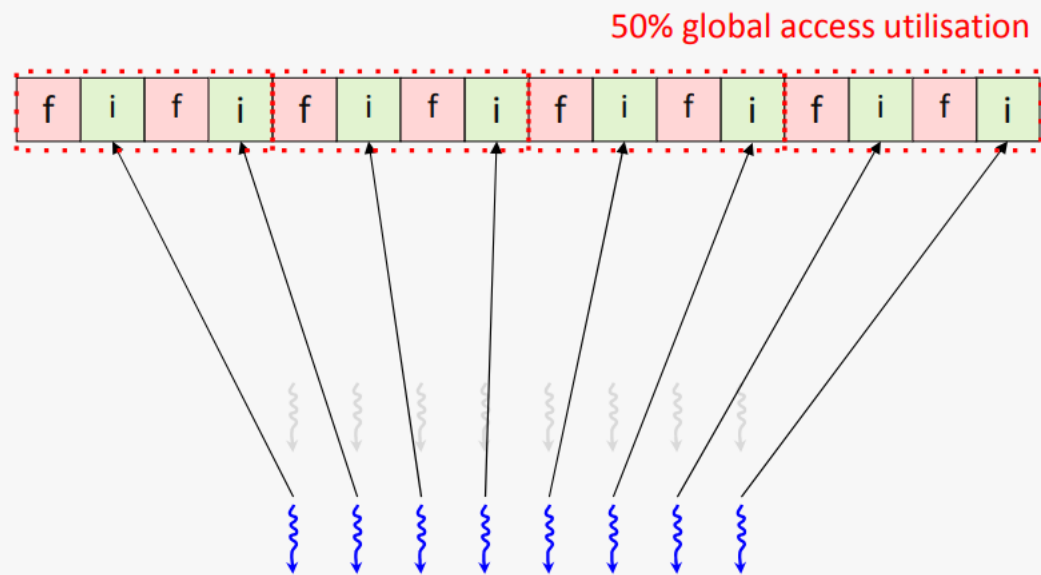
AOS VS SOA

```
struct str {  
    float f;  
    int i;  
};  
  
str data[size];  
  
...  
  
f(a[globalId].f);  
f(a[globalId].i);
```



AOS VS SOA


```
struct str {  
    float f;  
    int i;  
};  
  
str data[size];  
  
...  
  
f(a[globalId].f);  
f(a[globalId].i);
```



AOS VS SOA

```
struct str {  
    float fs[size];  
    int is[size];  
};
```

```
str data;
```




AOS VS SOA

```
struct str {  
    float fs[size];  
    int is[size];  
};
```

```
str data;
```


```
...
```

```
f(a[0].fs[globalId]);
```



AOS VS SOA

```
struct str {  
    float fs[size];  
    int is[size];  
};  
  
str data;  
  
...  
  
f(a[0].fs[globalId]);
```



AOS VS SOA


```
struct str {  
    float fs[size];  
    int is[size];  
};
```

```
str data;
```

```
...
```


```
f(a[0].fs[globalId]);
```

```
f(a[0].is[globalId]);
```




AOS VS SOA

```
struct str {  
    float fs[size];  
    int is[size];  
};  
  
str data;  
  
...  
  
f(a[0].fs[globalId]);  
f(a[0].is[globalId]);
```



COALESCED IMAGE CONVOLUTION PERFORMANCE



QUESTIONS

EXERCISE

`Code_Exercises/Exercise_16_Coalesced_Global_Memory/source`

Try inverting the dimensions when calculating the linear address in memory and measure the performance.